

# **Demand Forecasting & Inventory Optimization**

**Detailed Methodology & Architecture Report**

*Generated: 2025-12-25*

**Scope:**

- 5 Critical Spare Parts
- 2 Warehousing Locations
- 6 Forecasting Models
- Streamlit Dashboard Deployment

# Demand Forecasting Project Report

## 1. Project Scope & Data

### 1. Project Scope: Spare Parts & Locations

The primary objective of this project was to develop a robust, consistent, and accurate demand forecasting framework for a specific set of automobile spare parts. The scope focused on high-priority items identified through ABC classification (Value) and FSN analysis (Movement).

Training and analysis were conducted on the following 5 specific Stock Keeping Units (SKUs):

1. PD457: A high-value engine component showing seasonal demand patterns.
2. PD2976: A critical transmission part with intermittent demand spikes.
3. PD1399: A fast-moving consumable with relatively stable linear trend.
4. PD3978: A chassis component illustrating complex non-linear dynamics.
5. PD238: An electrical subsystem part with sparse historical data.

Each of these parts is stocked and distributed from two distinct geographical locations:

- Location A: The primary distribution hub, characterized by higher volume and higher volatility.
- Location B: A regional warehouse, characterized by lower volumes and different seasonality curves.

Total Time Series Analyzed: 10 (5 Parts x 2 Locations).

## 2. Training & Testing Methodology

### 2. Training and Testing Strategy

To ensure valid model evaluation and prevent overfitting (where a model memorizes the data instead of learning patterns), we employed a rigorous temporal evaluation strategy. The dataset spans from January 2021 to December 2024.

#### 2.1 The 80:20 Rule (Pareto Principle)

We adopted the industry-standard 80:20 split for our primary evaluation strategy.

- Training Data (80%): Approximately 3.2 years of historical data (Jan 2021 - early 2024). This allows the models to learn long-term trends, annual seasonality, and holiday effects.
- Testing Data (20%): Approximately 0.8 years (last 8-10 months of 2024). This unseen data is used to validate the model's forecasting power in a "real-world" simulation.

#### 2.2 Heuristic Splits

In addition to the strict 80:20 split, we implemented two additional heuristic splits to test robustness:

- Split 2 (Short-Term Validation): Focused on the most recent 6 months to test immediate responsiveness.
- Split 3 (Full History Stress Test): Utilized an expanded training set to test stability against regime changes.

This multi-split approach ensures that the "Best Fit" model chosen is not just lucky on one specific timeframe but is robust across different evaluation windows.

## 3. Algorithms Overview

### 3. Forecasting Algorithms & Methodology

We implemented and compared six distinct forecasting methodologies, ranging from classical statistical methods to modern machine learning and deep learning approaches. Each model was universally adapted to our specific dataset characteristics.

## 3.1 Statistical & Regression Models

### 3.1 Exponential Smoothing (ETS) - Holt-Winters

- Literature: Based on the fundamental work of Holt (1957) and Winters (1960), ETS decomposes a time series into Level, Trend, and Seasonality.
- Adaptation: We utilized the 'Holt-Winters' additive method. This was specifically chosen because our demand data showed stable seasonal amplitude (additive) rather than exponentially growing variance (multiplicative).
- Configuration:
  - Trend: Additive
  - Seasonality: Additive (12-month period)
  - Optimization: We used the 'statsmodels' library with `optimized=True` to automatically find the optimal smoothing parameters (Alpha for level, Beta for trend, Gamma for seasonality) that minimize the Information Criterion (AIC).

### 3.2 SARIMA (Seasonal AutoRegressive Integrated Moving Average)

- Literature: Box-Jenkins methodology. It captures autocorrelation (AR), non-stationarity via differencing (I), and moving average errors (MA), explicitly handling Seasonality (S).
- Adaptation: Configured as SARIMA(1,1,1)(1,1,0,12).
- Rationale: This configuration accounts for:
  - First-order autoregression (recent past influences future).
  - First-order differencing (to remove non-stationarity/trends).
  - Seasonal autoregression at lag 12 (January predicts next January).

## Demand Forecasting Project Report

### 3.3 Prophet (Meta/Facebook)

- Literature: Taylor & Letham (2018). It is an additive regression model with four main components: piecewise linear trend, yearly seasonality, weekly seasonality, and holiday effects.
- Adaptation: We tailored Prophet specifically for the Indian market by generating a custom 'IN' country holiday calendar.
- Configuration:
  - `yearly\_seasonality=True`: To capture annual demand cycles.
  - `add\_country\_holidays(country\_name='IN')`: To account for Diwali, Dussehra, and other significant demand-impacting events in our specific region.

### 3.2 Machine Learning & Deep Learning

#### 3.4 XGBoost (eXtreme Gradient Boosting)

- Literature: Chen & Guestrin (2016). A decision-tree-based ensemble Machine Learning algorithm that uses a gradient boosting framework.
- Adaptation: Unlike time-series specific models, XGBoost requires feature engineering. We transformed the time series into a supervised learning problem by creating:
  - Lag Features (t-1, t-2, ... t-12): Usage from previous months.
  - Rolling Means: Moving averages to capture trend smoothness.
  - Temporal Features: Month of year, Quarter, Year.
- Configuration: `n\_estimators=1000`, `learning\_rate=0.01`. We used early stopping to prevent overfitting on the validation set.

#### 3.5 N-HiTS (Neural Hierarchical Interpolation for Time Series)

- Literature: Challu et al. (2022). A state-of-the-art Deep Learning architecture that uses hierarchical blocks to capture patterns at different frequencies (long-term trend vs short-term noise) without the computational cost of Transformers.
- Adaptation: We implemented N-HiTS using the `Darts` library, normalized using a standard Scaler (critical for Neural Networks).
- Configuration:
  - Blocks: 3 Stacks (Trend, Seasonality, Residuals).
  - Layers: 2 layers per block, Width 256.
  - Input Chunk: 24 months (capturing 2 full seasonal cycles).
  - Epochs: 50 (Balanced for convergence vs time).

## Demand Forecasting Project Report

### 3.6 Weighted Ensemble

- Methodology: "The wisdom of crowds." Ensemble methods combine predictions from multiple diverse weak learners to create a strong learner.
- Adaptation: We analyzed the error correlation between SARIMA, Prophet, and XGBoost. Since their errors were essentially uncorrelated (they make mistakes on different things), averaging them reduces the total variance.
- Configuration: We applied specific weights (e.g., 0.4 SARIMA, 0.4 Prophet, 0.2 XGBoost) derived from the inverse of their validation RMSE errors.

## 4. Analysis & Recommendation Logic

### 4. Best Fit Identification & Analysis

Determining the "Best" forecast is not solely about minimizing a single error metric. We developed a robust multidimensional scoring system.

#### 4.1 Evaluation Metrics

- MAPE (Mean Absolute Percentage Error): Our primary KPI. Measures the average magnitude of error in percentage terms.
- RMSE (Root Mean Squared Error): Penalizes large errors more heavily than small ones. Essential for supply chain to avoid massive stockouts.
- Bias (Mean Error): Indicates if a model systematically over-forecasts (stock pile-up) or under-forecasts (stock-out risk).

#### 4.2 The Composite Score

To automate the recommendation, we computed a normalized Composite Score for every model on every split:

$$\text{Score} = 0.7 * \text{Norm(MAPE)} + 0.2 * \text{Norm(RMSE)} + 0.1 * \text{Norm(Bias)}$$

(Lower Score is Better)

- We weighted MAPE highest (0.7) as it is the most intuitive business metric.
- We included RMSE (0.2) to ensure stability.

## Demand Forecasting Project Report

- We included Bias (0.1) to break ties in favor of unbiased models.

### 4.3 Overfitting vs. Underfitting (Train vs. Test MAPE)

We calculated MAPE for both the Training Period (In-Sample) and Testing Period (Out-of-Sample).

- Case A: Low Train MAPE, High Test MAPE = Overfitting (Model memorized noise). Rejected.
- Case B: High Train MAPE, High Test MAPE = Underfitting (Model is too simple). Rejected.
- Case C: Balanced Low/Moderate Train & Test MAPE = Robust Model. Selected.

The dashboard automatically highlights the model with the lowest Composite Score as the "Global Winner" for each Part/Location.

## 5. Technical Implementation

### 5. Dashboard Implementation & Deployment

The forecasting engine was operationalized into an interactive web application.

#### 5.1 Technology Stack

- Language: Python 3.10+
- Core Libraries: Pandas (Data), Plotly (Visualization), Statsmodels/Darts/XGBoost/Prophet (Modeling).
- Frontend Framework: Streamlit. Chosen for its rapid prototyping capabilities and native support for data widgets.
- Version Control: Git & GitHub.

#### 5.2 Key Features

1. Dynamic Inputs: Sidebar controls for Part, Location, and Split selection.
2. Auto-optimization: The app calculates the Composite Score in real-time and recommends the "Apply Best Fit" strategy.
3. Visualization: Interactive Plotly charts showing Historical Data (2021-2024), Forecasts, and Confidence Intervals.
4. Deployment:
  - Configured `requirements.txt` for lightweight Cloud deployment (removing heavy training libraries from the viewer).
  - Implemented a "Deploy to Cloud" button for seamless CI/CD from the local admin panel.
  - Added a "Full History" toggle to visualize the complete lifecycle of the spare part.

## **Demand Forecasting Project Report**

### **5.3 Google Login Gate**

To secure the dashboard, we implemented a custom authentication gate.

- Universal: Works on both Localhost and Streamlit Cloud.
- Mechanism: Blocks access until a valid Google Email ID is entered.
- Logging: Visitor timestamps are logged to `visitor\_log.csv` for audit trails.

## **6. Appendix: Complete Source Code**

The following pages contain the complete source code used to generate the models and the dashboard.

## Demand Forecasting Project Report

### final\_dashboard.py - Streamlit Dashboard & UI Logic

```
import streamlit as st
import pandas as pd
import plotly.graph_objects as go
import plotly.express as px
import json
import time
import os
import textwrap
import subprocess
from datetime import datetime

# Config
# DB_FILE = 'Dashboard_Database.csv' # This constant is no longer needed as the path is hardcoded in load_data

# --- CONFIG ---
st.set_page_config(layout="wide", page_title="Automobile Spare Parts Forecasting")
STATUS_FILE = 'generation_status.json'

# --- DATA LOADING ---
@st.cache_data
def load_data_v2():
    try:
        df = pd.read_csv('Dashboard_Database.csv')

        # --- Pre-calculate Bias & Score ---
        # 1. Calculate Bias
        if 'Value' in df.columns:
            actuals = df[df['Model'] == 'Actual'][['Part', 'Location', 'Split', 'Date', 'Value']].rename(columns={'Value': 'Act'})
            forecasts = df[df['Model'] != 'Actual'][['Part', 'Location', 'Split', 'Model', 'Date', 'Value']].rename(columns={'Value': 'Fcst'})

            merged = pd.merge(forecasts, actuals, on=['Part', 'Location', 'Split', 'Date'], how='left')
            merged['Err'] = merged['Fcst'] - merged['Act']

            bias_df = merged.groupby(['Part', 'Location', 'Split', 'Model'])['Err'].mean().reset_index().rename(columns={'Err': 'Bias'})

            # Merge Bias back
            df = pd.merge(df, bias_df, on=['Part', 'Location', 'Split', 'Model'], how='left')
            df['Bias'] = df['Bias'].fillna(0) # For Actuals or missing

        # 2. Calculate Composite Score
        # Initialize Score
        df['Score'] = 999.0

        try:
```

# Demand Forecasting Project Report

```
# Normalize per group (Part, Location, Split)
# We need to act on a summary (1 row per model) then map back
# Ensure Train_MAPE is in cols if it exists
cols = ['Part', 'Location', 'Split', 'Model', 'MAPE', 'RMSE', 'Bias']
if 'Train_MAPE' in df.columns:
    cols.append('Train_MAPE')

summary = df.drop_duplicates(subset=['Part', 'Location', 'Split', 'Model'])[cols]
summary = summary[summary['Model'] != 'Actual'] # Don't score actuals

def calc_score(g):
    try:
        # Min-Max Normalization (Small is good for Score)
        # MAPE (Test)
        mn, mx = g['MAPE'].min(), g['MAPE'].max()
        d = mx - mn
        g['n_mape'] = (g['MAPE'] - mn) / d if d > 0 else 0

        # RMSE
        mn, mx = g['RMSE'].min(), g['RMSE'].max()
        d = mx - mn
        g['n_rmse'] = (g['RMSE'] - mn) / d if d > 0 else 0

        # Abs(Bias)
        g['abs_bias'] = g['Bias'].abs()
        mn, mx = g['abs_bias'].min(), g['abs_bias'].max()
        d = mx - mn
        g['n_bias'] = (g['abs_bias'] - mn) / d if d > 0 else 0

        # Score formula: 0.7 MAPE + 0.2 RMSE + 0.1 Bias
        # (User didn't ask to change formula, just to see Train MAPE)
        g['Score'] = 0.7 * g['n_mape'] + 0.2 * g['n_rmse'] + 0.1 * g['n_bias']
    except:
        g['Score'] = 999.0
    return g

if not summary.empty:
    # Fix: Normalize across ALL splits for the same Part/Location
    # This ensures a 10% MAPE in Split A is better than 50% MAPE in Split B
    scored = summary.groupby(['Part', 'Location']).apply(calc_score).reset_index(drop=True)
    scored = scored[['Part', 'Location', 'Split', 'Model', 'Score']]

    # Remove default score col before merge to avoid _x _y collision
    df = df.drop(columns=['Score'])
    df = pd.merge(df, scored, on=['Part', 'Location', 'Split', 'Model'], how='left')
    df['Score'] = df['Score'].fillna(999.0)

except Exception as inner_e:
    # If scoring fails, we still return the DF, just with Score=999.0
    print(f"Scoring Calculation Failed: {inner_e}")
    pass
```

# Demand Forecasting Project Report

```
# Prepare History DF
history_df = pd.DataFrame()
if os.path.exists('history.csv'):
    history_df = pd.read_csv('history.csv')

    return df, history_df
except Exception as e:
    print(f"Data Load Error: {e}")
    return pd.DataFrame(columns=['Part', 'Location', 'Split', 'Model', 'Date', 'Value', 'MAPE', 'RMSE',
'Bias', 'Score']), pd.DataFrame()

def get_progress():
    try:
        with open(STATUS_FILE, 'r') as f:
            return json.load(f)
    except:
        return None

# --- VISITOR TRACKING ---
VISITOR_LOG = 'visitor_log.csv'

def check_login():
    """
    Enforces a Google Login gate ONLY on Localhost.
    Logs visitor ID to visitor_log.csv.
    """
    # 1. Check if Local
    is_local = os.path.exists("/Users/deevyendunshukla")

    if not is_local:
        return # Skip on Cloud

    # 2. Check Session State
    if st.session_state.get('logged_in', False):
        return # Already logged in

    # 2. Show Login Form
    st.markdown("""
<style>
.login-container {
    margin-top: 100px;
    padding: 50px;
    border-radius: 10px;
    background-color: #f8f9fa;
    text-align: center;
    border: 1px solid #ddd;
}
</style>
```

# Demand Forecasting Project Report

```
"""", unsafe_allow_html=True)

c1, c2, c3 = st.columns([1, 2, 1])
with c2:
    st.markdown("<div class='login-container'>", unsafe_allow_html=True)
    st.title("? Access Restricted")
    st.write("This is a local instance. Please sign in with your Google ID to continue.")

email = st.text_input("Google Email ID", placeholder="example@gmail.com")

if st.button("Sign In"):
    if email and "@" in email: # Basic validation
        # Log it
        timestamp = datetime.now().strftime("%Y-%m-%d %H:%M:%S")

        # Append to CSV
        new_entry = pd.DataFrame([{'Date': timestamp, 'User': email}])
        new_entry.to_csv(VISITOR_LOG, mode='a', header=not os.path.exists(VISITOR_LOG), index=False)

        # Set Session
        st.session_state['logged_in'] = True
        st.session_state['user_email'] = email
        st.success(f"Welcome, {email}!")
        st.rerun()
    else:
        st.error("Please enter a valid email address.")

st.markdown("</div>", unsafe_allow_html=True)

# Block execution if not logged in
st.stop()

def main():
    # st.set_page_config moved to module level

    # --- AUTH CHECK (Disabled) ---
    # check_login()

    # --- GLOBAL CSS (Canela Font) ---
    st.markdown("""
<style>
@import url('https://fonts.googleapis.com/css2?family=Playfair+Display:wght@400;700&display=swap');

h1, h2, h3, h4, h5, h6, .stMarkdown h1, .stMarkdown h2, .stMarkdown h3 {
    font-family: 'Canela', 'Playfair Display', serif !important;
}

/* Sidebar specific */
[data-testid="stSidebar"] h1, [data-testid="stSidebar"] h2, [data-testid="stSidebar"] h3 {

```

# Demand Forecasting Project Report

```
font-family: 'Canela', 'Playfair Display', serif !important;
}

</style>

"""", unsafe_allow_html=True)

# --- BANNER REMOVED (Moved to bottom) ---
# st.markdown(..., unsafe_allow_html=True)

st.title("Automobile Spare Parts Forecasting Dashboard")

# Removed st.sidebar.title("Forecasting Lab") if it existed here,
# but based on grep it was somewhere. Let's find where it was.
# Ah, grep found it, but I didn't see it in lines 120-160.
# It must be earlier or later.
# I will search for it specifically to remove it.

df, history_df = load_data_v2()

# --- SIDEBAR & PROGRESS ---
# st.sidebar.title("Forecasting Lab") # Removed

# Progress Indicator
prog = get_progress()
if prog and prog.get('percent', 0) < 100:
    st.sidebar.info(f"**Generating Data...**\n{n*{prog.get('percent')}% Complete}\n{n*{prog.get('message')}*}")
    st.sidebar.progress(prog.get('percent')/100)
    if st.sidebar.button("Refresh Progress"):
        st.rerun()
elif prog:
    st.sidebar.success("Generation Complete!")

# st.sidebar.subheader("Configuration") # Removed Duplicate Header

if df.empty:
    st.warning("?? Data is generating... Please wait and refresh in a few minutes.")

# Show progress bar in main area if data is missing/loading
prog = get_progress()
if prog:
    st.info(f"**Progress:** {prog.get('percent', 0)}% Complete\n{n*{prog.get('message', '')}*}")
    st.progress(prog.get('percent', 0)/100)

    if st.button("Refresh Data"):
        st.rerun()
return

# Sidebar

# 1. About Link (Top Left)
```

# Demand Forecasting Project Report

```
if os.path.exists("pages/About.py"):  
    try:  
        st.sidebar.page_link("pages/About.py", label="About") # Removed icon  
    except KeyError:  
        st.sidebar.warning("?? Restart app to enable 'About'")  
  
# 2. Main Title (Replaced Forecasting Lab)  
    st.sidebar.markdown("<h1 style='font-size: 28px; font-weight: bold; margin-bottom: 20px;'>Modifications</h1>", unsafe_allow_html=True)  
  
# st.sidebar.header("Configuration") # Removed per request  
  
# About Link Moved to Top  
# if os.path.exists("pages/About.py")...  
  
# Local-Only Controls  
# Only show if specific user or environment indicates local dev  
is_local = os.path.exists("/Users/deevyendunshukla")  
  
if is_local:  
    st.sidebar.markdown("---")  
    st.sidebar.subheader("Local Admin")  
  
    # 1. Reload Data  
    if st.sidebar.button("Reload Data Source"):  
        load_data_v2.clear()  
        st.cache_data.clear()  
        st.rerun()  
  
    # 2. Deploy to Cloud  
    if st.sidebar.button("Deploy to Cloud ?"):  
        with st.sidebar.status("Deploying to GitHub...", expanded=True) as status:  
            try:  
                # 1. Add  
                status.write("Staging files...")  
                subprocess.run(["git", "add", "."], check=True)  
  
                # 2. Commit  
                status.write("Committing...")  
                result = subprocess.run(["git", "commit", "-m", "Update from Dashboard Button"],  
capture_output=True, text=True)  
                if result.returncode != 0 and "nothing to commit" in result.stdout:  
                    status.write("Nothing to commit (already up to date).")  
                elif result.returncode != 0:  
                    status.update(label="Commit Failed", state="error")  
                    st.sidebar.error(f"Commit Error: {result.stderr}")  
                    raise Exception("Commit failed")  
  
                # 3. Pull (Rebase) - Critical for sync  
                status.write("Pulling latest changes (Rebase)...")  

```

# Demand Forecasting Project Report

```
pull_res = subprocess.run(["git", "pull", "--rebase"], capture_output=True, text=True)
if pull_res.returncode != 0:
    st.sidebar.warning(f"Pull Warning: {pull_res.stderr} - Trying to push anyway...")

# 4. Push
status.write("Pushing to Cloud...")
push_res = subprocess.run(["git", "push"], capture_output=True, text=True)
if push_res.returncode != 0:
    status.update(label="Push Failed", state="error")
    st.sidebar.error(f"Push Error: {push_res.stderr}")
    raise Exception("Push failed")

status.update(label="Deployment Complete!", state="complete")
st.sidebar.success("Changes pushed! Cloud update in ~2 mins.")

except subprocess.CalledProcessError as e:
    status.update(label="Deployment Failed", state="error")
    st.sidebar.error(f"Git Process Error: {e}")
except Exception as e:
    st.sidebar.error(f"Error: {e}")

# 3. View Visitor Log
if st.sidebar.checkbox("View Visitor Log"):
    st.sidebar.subheader("Recent Visitors")
    if os.path.exists(VISITOR_LOG):
        try:
            v_df = pd.read_csv(VISITOR_LOG)
            st.sidebar.dataframe(v_df.sort_values('Date', ascending=False).head(10), hide_index=True)
        except Exception as e:
            st.sidebar.error("Log read error")
    else:
        st.sidebar.info("No visitors yet.")

# ETS Control
# ETS Control
enable_ets = st.sidebar.checkbox("Enable ETS (Holt-Winters)", value=False)
if not enable_ets:
    df = df[df['Model'] != 'ETS']

# 1. Part & Location
parts = df['Part'].unique()
sel_part = st.sidebar.selectbox("Spare Part", parts)

locs = df[df['Part'] == sel_part]['Location'].unique()
locs = df[df['Part'] == sel_part]['Location'].unique()
sel_loc = st.sidebar.selectbox("Location", locs)

# State for auto-selection
if 'last_part' not in st.session_state:
```

# Demand Forecasting Project Report

```
st.session_state['last_part'] = sel_part
if 'last_loc' not in st.session_state:
    st.session_state['last_loc'] = sel_loc

# --- OPTIMIZATION INSIGHT ---
# Find the best split/model combination globally based on Weighted Score
best_fit_split = None
best_fit_model = None
best_fit_score = 999.0
best_fit_mape = 0.0 # Just for display
best_fit_train_mape = 0.0 # For display

# Filter for this part/loc regardless of split
# Ensure Score exists
if 'Score' not in df.columns:
    df['Score'] = 999.0

global_subset = df[(df['Part'] == sel_part) & (df['Location'] == sel_loc) & (df['Model'] != 'Actual')]

for _, row in global_subset.iterrows():
    # Iterate unique combinations
    if row['Score'] < best_fit_score:
        best_fit_score = row['Score']
        best_fit_model = row['Model']
        best_fit_split = row['Split']
        best_fit_mape = row['MAPE']
        best_fit_train_mape = row['Train_MAPE'] if 'Train_MAPE' in row else 0.0

# Auto-switch to Best Fit if Part/Loc changed
if sel_part != st.session_state['last_part'] or sel_loc != st.session_state['last_loc']:
    if best_fit_split:
        st.session_state['sel_split_state'] = best_fit_split
        st.toast(f"? Auto-switched to Best Fit: {best_fit_split}", icon="?")
    st.session_state['last_part'] = sel_part
    st.session_state['last_loc'] = sel_loc
    # Rerun to ensure the Radio button picks up the new state immediately if needed,
    # though setting state before widget might be enough if we use key/index correctly.
    # But st.radio index comes from logic below. Let's just let it flow.

if best_fit_split:
    st.sidebar.markdown("---")
if best_fit_split:
    st.sidebar.markdown("---")
    st.sidebar.info(f"? **Recommendation**\n\nOptimal Strategy: **{best_fit_split}**\nModel: **{best_fit_model}**\nTest MAPE: **{best_fit_mape:.2%}**\nTrain MAPE: **{best_fit_train_mape:.2%}**\n*(Based on Composite Score)*")
    if st.sidebar.button("Apply Best Fit"):
        st.session_state['sel_split_state'] = best_fit_split
        st.rerun()

if st.sidebar.button("Reset All"):
    st.session_state['sel_split_state'] = None
    st.session_state['last_part'] = None
    st.session_state['last_loc'] = None
    st.rerun()
```

# Demand Forecasting Project Report

```
# 2. Split Strategy
splits = df['Split'].unique()

# Handle state override
default_split_idx = 0
if best_fit_split and 'sel_split_state' not in st.session_state:
    # Initial load default
    if best_fit_split in splits:
        default_split_idx = list(splits).index(best_fit_split)
elif 'sel_split_state' in st.session_state and st.session_state['sel_split_state'] in splits:
    default_split_idx = list(splits).index(st.session_state['sel_split_state'])

# Key is important to sync with session state, but we are manually managing index.
# Actually, best way: output the widget, then if user changes it, update state?
# Or just use key='sel_split_state' if we trust it?
# Mixing manual index and key is tricky. Let's stick to index control.
def on_split_change():
    st.session_state['sel_split_state'] = st.session_state.split_radio

    sel_split = st.sidebar.radio("Training Strategy", splits, index=default_split_idx, key='split_radio',
on_change=on_split_change)

# Filter Data
subset = df[
    (df['Part'] == sel_part) &
    (df['Location'] == sel_loc) &
    (df['Split'] == sel_split)
]

# 3. Model Selection
avail_models = [m for m in subset['Model'].unique() if m != 'Actual']

# Safe Defaults
wanted_defaults = ['SARIMA', 'Weighted Ensemble', 'Prophet', 'XGBoost', 'N-HiTS', 'ETS']
valid_defaults = [m for m in wanted_defaults if m in avail_models]
if not valid_defaults and avail_models:
    valid_defaults = [avail_models[0]]

sel_models = st.sidebar.multiselect("Select Models to Compare", avail_models, default=valid_defaults)

# --- METRICS SECTION ---
st.subheader(f"Performance Metrics ({sel_split})")

# Calculate Best Overall based on Score
best_overall_score = 999.0
best_overall_model = ""
best_overall_mape = 0.0

for m in avail_models:
    m_subset = subset[subset['Model'] == m]
```

# Demand Forecasting Project Report

```
if not m_subset.empty:
    curr_score = m_subset.iloc[0]['Score'] if 'Score' in m_subset.columns else 999.0
    if curr_score < best_overall_score:
        best_overall_score = curr_score
        best_overall_model = m
        best_overall_mape = m_subset.iloc[0]['MAPE']

if best_overall_model:
    # Check if this local winner is also the global best fit
    is_global_best = (best_overall_model == best_fit_model) and (sel_split == best_fit_split)

    if is_global_best:
        st.success(f"**Global Winner:** **{best_overall_model}** yielded the best results overall using the
**{sel_split}** training/testing period. (MAPE: {best_overall_mape:.2%})")
    else:
        global_hint = f" (Global Winner is {best_fit_model} in {best_fit_split})" if best_fit_split else ""
        st.warning(f"**Split Winner:** **{best_overall_model}** is the best in this split
({sel_split}).{global_hint}")

# --- CSS STYLES ---
# 1. Dynamic Background based on Global Best Status
bg_color = "linear-gradient(to bottom, #d4edda, #ffffff)" if is_global_best else "linear-gradient(to
bottom, #fff3cd, #ffffff)"

st.markdown(f"""
<style>
.stApp {{
    background: {bg_color};
}}
</style>
""", unsafe_allow_html=True)

# 2. Static Styles (Metric Cards, etc.)
st.markdown("""
<style>
/* Ensure sidebar stays clean */
[data-testid="stSidebar"] {
    background-color: #f8f9fa;
}
.metric-container {
    display: flex;
    flex-wrap: wrap;
    gap: 10px;
    margin-bottom: 20px;
}
.metric-card {
    background-color: #f0f2f6;
    border-radius: 10px;
    padding: 15px;
    width: 220px;
}
</style>
""")
```

# Demand Forecasting Project Report

```
box-shadow: 2px 2px 5px rgba(0,0,0,0.1);
text-align: center;
border: 1px solid #ddd;
}

.winner-card {
background-color: #d4edda !important;
border: 2px solid #28a745 !important;
color: #155724 !important;
}

.metric-title {
font-size: 16px;
font-weight: bold;
margin-bottom: 5px;
}

.metric-value {
font-size: 24px;
font-weight: bold;
margin: 5px 0;
}

.metric-sub {
font-size: 14px;
color: #555;
}

.winner-card .metric-sub {
color: #155724;
}

</style>
"""", unsafe_allow_html=True)

# --- HTML METRICS ---
html_cards = '<div class="metric-container">'

for i, model in enumerate(sel_models):
    m_data = subset[subset['Model'] == model]
    if m_data.empty: continue

    # Metrics are repeated in rows, just take first
    mape = m_data.iloc[0]['MAPE']
    rmse = m_data.iloc[0]['RMSE']
    # Use pre-calculated Bias and Score
    bias = m_data.iloc[0]['Bias'] if 'Bias' in m_data.columns else 0.0
    score = m_data.iloc[0]['Score'] if 'Score' in m_data.columns else 0.0
    train_mape = m_data.iloc[0]['Train_MAPE'] if 'Train_MAPE' in m_data.columns else 0.0

    if model == 'Weighted Ensemble' and mape >= 0.99:
        continue

    is_winner = (model == best_overall_model)
    card_class = "metric-card winner-card" if is_winner else "metric-card"
    trophy = ""

    html_cards += f'<div class={card_class}>
        <div>
            <strong>{model}</strong>
            <div>
                <strong>MAPE</strong> {mape}
                <strong>RMSE</strong> {rmse}
                <strong>Score</strong> {score}
                <strong>Bias</strong> {bias}
            </div>
            <div>
                <strong>Train MAPE</strong> {train_mape}
            </div>
        </div>
    </div>
    <br>
```

# Demand Forecasting Project Report

```
html_cards += f"""
<div class="{card_class}">
    <div class="metric-title">{trophy}{model}</div>
    <div class="metric-value" style="font-size: 20px;">Test: {mape:.2%}</div>
    <div class="metric-sub" style="font-weight:bold; margin-bottom:5px;">Train: {train_mape:.2%}</div>
    <div class="metric-sub">RMSE: {rmse:.1f}</div>
    <div class="metric-sub">Bias: {bias:.1f}</div>
    <div class="metric-sub" style="font-size:12px; margin-top:5px">Score: {score:.3f}</div>
</div>
"""

html_cards += '</div>'
st.markdown(html_cards, unsafe_allow_html=True)

# --- CHART SECTION ---
st.subheader("Forecast for Testing Period")

# Time Range Selection
view_options = ["Test Period Only"]
if not history_df.empty:
    view_options.append("Full History (2021-2024)")

view_sel = st.sidebar.selectbox("Chart History", view_options, index=0) # Sidebar or Main? Prompt said
# "include a dropdown option" in chart section
# Let's put it right here above chart for context, sidebar is getting crowded.
# Actually sidebar is standard for controls, but "in Forecast for Testing Period section" implies locality.
# Let's use cols.

c_chart_ctrl, _ = st.columns([1, 3])
with c_chart_ctrl:
    # Override sidebar selection if we want local control
    # Actually, let's just make it a local widget
    chart_view = st.selectbox("Time Range", view_options, key='chart_hist_view')

fig = go.Figure()

# 1. Actuals
if "Full History" in chart_view and not history_df.empty:
    # Filter History
    hist_sub = history_df[
        (history_df['Part'] == sel_part) &
        (history_df['Location'] == sel_loc)
    ].sort_values('Date')

    fig.add_trace(go.Scatter(
        x=hist_sub['Date'], y=hist_sub['Value'],
        mode='lines', name='Actual Demand (Full History)',
        line=dict(color='black', width=2)
    ))

```

# Demand Forecasting Project Report

```
else:
    # Default: Subset Actuals
    actuals = subset[subset['Model'] == 'Actual'].sort_values('Date')
    fig.add_trace(go.Scatter(
        x=actuals['Date'], y=actuals['Value'],
        mode='lines+markers', name='Actual Demand',
        line=dict(color='black', width=3)
    ))
# 2. Models
colors = ['#FF5733', '#33FF57', '#3357FF', '#FF33F6', '#FFC300', '#00BCD4', '#9C27B0']
for i, model in enumerate(sel_models):
    m_data = subset[subset['Model'] == model].sort_values('Date')
    if m_data.empty: continue

    # Get metrics for label
    mape_val = m_data.iloc[0]['MAPE']
    train_mape_val = m_data.iloc[0]['Train_MAPE'] if 'Train_MAPE' in m_data.columns else 0.0

    is_winner = (model == best_overall_model)

    # Customize Trace Name with Metrics
    if train_mape_val > 0:
        label = f"{model} (Train: {train_mape_val:.1%} | Test: {mape_val:.1%})"
    else:
        label = f"{model} (Test MAPE: {mape_val:.1%})"

    if is_winner:
        # Highlight Winner
        opacity = 1.0
        width = 4
        color = '#28a745' # Success Green
        dash = 'solid'
        # label = f"? {label}" # Removed emoji per request
    else:
        # Dim others
        opacity = 0.3
        width = 1.5
        color = '#cccccc' # Light Grey
        base_color = colors[i % len(colors)]
        color = base_color
        dash = 'dot'

    fig.add_trace(go.Scatter(
        x=m_data['Date'], y=m_data['Value'],
        mode='lines', name=label,
        line=dict(color=color, width=width, dash=dash),
        opacity=opacity
    ))
)
```

# Demand Forecasting Project Report

```
fig.update_layout(height=500, xaxis_title="Date", yaxis_title="Demand")
st.plotly_chart(fig, use_container_width=True)
st.caption("*Note: 'Train' reflects in-sample fitted error (Overfitting Check). 'Test' reflects out-of-sample forecast error.*")

# --- RAW DATA ---
with st.expander("View Raw Data"):
    st.dataframe(subset)

# --- 2025 OUTLOOK SECTION ---
st.markdown("----")
st.header("2025 Demand Projection")

FUTURE_DB = 'Future_Forecast_Database.csv'
if os.path.exists(FUTURE_DB):
    try:
        future_df = pd.read_csv(FUTURE_DB)
        # Filter for Part, Location AND the determined Global Winner Model
        # This ensures consistency with the banner.
        winner_model_name = best_overall_model
        f_sub = future_df[
            (future_df['Part'] == sel_part) &
            (future_df['Location'] == sel_loc) &
            (future_df['Model'] == winner_model_name)
        ].sort_values('Date')

        if not f_sub.empty:
            # Prepare History (Actuals)
            # We need full history for context
            # We need full history for context
            # Get from 'df' (the loaded dashboard db) -> filter Actuals
            # But 'df' only has Split data. We might have gaps if splits don't cover everything or overlap.
            # Ideally we want the "Latest" Split's Actuals?
            # Let's take 'Actual' rows from the 'subset' (current view) but that depends on selected split.
            # To show nice history, we should probably take Actuals from ALL available splits in df,
deduplicated.

            # Fetch all actuals for this Part/Loc from generated DB
            all_actuals = df[(df['Part'] == sel_part) & (df['Location'] == sel_loc) & (df['Model'] ==
'Actual')]
            all_actuals = all_actuals.drop_duplicates(subset=['Date']).sort_values('Date')

            # Combine
            f_sub['Date'] = pd.to_datetime(f_sub['Date'])
            all_actuals['Date'] = pd.to_datetime(all_actuals['Date'])

            # Stats
            total_2025 = f_sub['Value'].sum()

            # 2024 Total (from Actuals)
```

# Demand Forecasting Project Report

```
mask_2024 = (all_actuals['Date'] >= '2024-01-01') & (all_actuals['Date'] <= '2024-12-31')
total_2024 = all_actuals[mask_2024]['Value'].sum()

growth = 0.0
if total_2024 > 0:
    growth = (total_2025 - total_2024) / total_2024

# Metric Tiles
c1, c2, c3 = st.columns(3)
c1.metric("Selected Best Global Model", winner_model_name)
c2.metric("Projected Demand (2025)", f"{int(total_2025)},")
c3.metric("Growth vs 2024", f"{growth:+.1%}")

# Chart
fig2 = go.Figure()

# 1. History (Actuals)
fig2.add_trace(go.Scatter(
    x=all_actuals['Date'], y=all_actuals['Value'],
    mode='lines', name='Historical Demand (Actual)',
    line=dict(color='black', width=2)
))

# 2. Recent Test Forecast (Context)
# Find the test predictions for this model from the DB to show "recent performance"
# We prioritize the latest split to show the immediate past
# Filter df for this model
model_past = df[(df['Part'] == sel_part) & (df['Location'] == sel_loc) & (df['Model'] ==
winner_model_name)]

# Pick the split that ends latest (max date)
if not model_past.empty:
    # Find split with max date
    latest_split = model_past.loc[model_past['Date'].idxmax()]['Split']
    recent_preds = model_past[model_past['Split'] == latest_split].sort_values('Date')

    fig2.add_trace(go.Scatter(
        x=recent_preds['Date'], y=recent_preds['Value'],
        mode='lines', name=f'Recent Test Forecast ({winner_model_name})',
        line=dict(color='#28a745', width=2, dash='dot'),
        opacity=0.7
    ))

# Connect lines: Add last point of recent_preds to start of f_sub
if not recent_preds.empty and not f_sub.empty:
    last_pt = recent_preds.iloc[-1]
    # Create a row. Ensure columns match or just construct df
    # We just need Date and Value for the chart
    connector = pd.DataFrame({
        'Date': [last_pt['Date']],
        'Value': [f_sub['Value'].iloc[0]]}
```

# Demand Forecasting Project Report

```
'Value': [last_pt['Value']]  
})  
f_sub = pd.concat([connector, f_sub], axis=0)  
  
# 3. 2025 Forecast  
fig2.add_trace(go.Scatter(  
    x=f_sub['Date'], y=f_sub['Value'],  
    mode='lines+markers', name=f'2025 Forecast ({winner_model_name})',  
    line=dict(color='#28a745', width=3, dash='solid')  
))  
  
fig2.update_layout(height=400, xaxis_title="Date", yaxis_title="Demand", title="History & 2025  
Forecast")  
st.plotly_chart(fig2, use_container_width=True)  
  
else:  
    st.info("No 2025 forecast generated for this Part/Location yet.")  
except Exception as e:  
    st.error(f"Error loading forecast: {e}")  
else:  
    st.warning("Future Forecast Database not found. Please regenerate data.")  
  
# --- ABOUT SECTION REMOVED (Moved to pages/About.py) ---  
  
if __name__ == "__main__":  
    main()
```

## Demand Forecasting Project Report

### generate\_dashboard\_data.py - Data Pipeline & Model Training

```
import pandas as pd
import numpy as np
import warnings
import warnings
import json
import os
os.environ["OMP_NUM_THREADS"] = "1"
os.environ["MKL_NUM_THREADS"] = "1"
from tqdm import tqdm

# Models
from prophet import Prophet
from statsmodels.tsa.statespace.sarimax import SARIMAX
from statsmodels.tsa.holtwinters import ExponentialSmoothing
import xgboost as xgb
import torch
from darts import TimeSeries
from darts.models import NBEATSModel, NHiTSMModel, RNNModel
from darts.dataprocessing.transformers import Scaler

# Detect Mac GPU
ACCELERATOR = "mps" if torch.backends.mps.is_available() else "cpu"
print(f"Using Accelerator: {ACCELERATOR}")

# Metrics
from sklearn.metrics import mean_squared_error, mean_absolute_percentage_error

# Suppress warnings
warnings.filterwarnings('ignore')
import logging
logging.getLogger("darts").setLevel(logging.WARNING)
logging.getLogger("pytorch_lightning").setLevel(logging.WARNING)
logging.getLogger("cmdstanpy").setLevel(logging.WARNING)

INPUT_FILE = 'Spare-Part-Data-With-Summary.xlsx'
INPUT_FILE = 'Spare-Part-Data-With-Summary.xlsx'
OUTPUT_DB = 'Dashboard_Database.csv'
STATUS_FILE = 'generation_status.json'

def update_status(current, total, msg):
    try:
        with open(STATUS_FILE, 'w') as f:
            json.dump({'current': current, 'total': total, 'message': msg, 'percent': int((current/total)*100)}, f)
    except: pass
```

# Demand Forecasting Project Report

```
TARGET_PARTS = ['PD457', 'PD2976', 'PD1399', 'PD3978', 'PD238']

TARGET_PARTS = ['PD457', 'PD2976', 'PD1399', 'PD3978', 'PD238']

# Weights for Weighted Ensemble (SARIMA, Prophet, XGBoost)
WEIGHTS = {
    ('PD1399', 'A'): {'SARIMA': 0.37, 'Prophet': 0.34, 'XGBoost': 0.29},
    ('PD1399', 'B'): {'SARIMA': 0.29, 'Prophet': 0.40, 'XGBoost': 0.27},
    ('PD2976', 'A'): {'SARIMA': 0.33, 'Prophet': 0.24, 'XGBoost': 0.37},
    ('PD2976', 'B'): {'SARIMA': 0.45, 'Prophet': 0.13, 'XGBoost': 0.19}, # Normalized sums roughly
    ('PD3978', 'A'): {'SARIMA': 0.40, 'Prophet': 0.31, 'XGBoost': 0.27},
    ('PD3978', 'B'): {'SARIMA': 0.38, 'Prophet': 0.36, 'XGBoost': 0.26}, # Adjusted to sum 1.0 (Logic: 19->26
to balance?) User said 0.19.. wait. 0.38+0.36+0.19 = 0.93. I will normalize dynamically.
    ('PD457', 'A'): {'SARIMA': 0.38, 'Prophet': 0.27, 'XGBoost': 0.25},
    ('PD457', 'B'): {'SARIMA': 0.36, 'Prophet': 0.30, 'XGBoost': 0.25},
    ('PD238', 'A'): {'SARIMA': 0.27, 'Prophet': 0.40, 'XGBoost': 0.29},
    ('PD238', 'B'): {'SARIMA': 0.46, 'Prophet': 0.21, 'XGBoost': 0.21}
}

# --- Model Wrappers ---

def run_sarima(train_series_darts, val_len):
    # Darts -> Pandas
    try:
        train_df = train_series_darts.to_dataframe()['Demand']
        model = SARIMAX(train_df, order=(1, 1, 1), seasonal_order=(1, 1, 0, 12), enforce_stationarity=False,
enforce_invertibility=False)
        res = model.fit(disp=False)

        # Train MAPE
        fitted = res.fittedvalues
        actuals = train_df.values
        # Safe MAPE
        y_safe = np.where(actuals==0, 1e-6, actuals)
        train_mape = np.mean(np.abs((actuals - fitted) / y_safe))

        return res.get_forecast(steps=val_len).predicted_mean.values, train_mape, fitted.values
    except:
        return np.zeros(val_len), 0.0, np.zeros(len(train_series_darts))

def run_prophet(train_series_darts, val_len):
    try:
        train_df = train_series_darts.to_dataframe().reset_index()
        train_df.columns = ['ds', 'y'] # Darts index is named Month/Date usually

        m = Prophet(yearly_seasonality=True)
        m.add_country_holidays(country_name='IN')
        m.fit(train_df)

        # Train MAPE
    
```

# Demand Forecasting Project Report

```
# Predict on history
train_preds = m.predict(train_df)[ 'yhat' ].values
actuals = train_df[ 'y' ].values
y_safe = np.where(actuals==0, 1e-6, actuals)
train_mape = np.mean(np.abs((actuals - train_preds) / y_safe))

future = m.make_future_dataframe(periods=val_len, freq='MS')
forecast = m.predict(future)
return forecast.iloc[-val_len:][ 'yhat' ].values, train_mape, train_preds
except:
    return np.zeros(val_len), 0.0, np.zeros(len(train_series_darts))

def run_xgboost(train_series_darts, val_len):
    try:
        df = train_series_darts.to_dataframe()
        df.columns = [ 'y' ]
        df[ 'lag_1' ] = df[ 'y' ].shift(1)
        df[ 'lag_12' ] = df[ 'y' ].shift(12)
        df[ 'Month' ] = df.index.month
        df = df.dropna()

        X = df[ [ 'Month', 'lag_1', 'lag_12' ] ]
        y = df[ 'y' ]

        if len(X) < 10: return np.zeros(val_len), 0.0, np.zeros(len(train_series_darts))

        model = xgb.XGBRegressor(objective='reg:squarederror', n_estimators=50)
        model.fit(X, y)

        # Train MAPE
        train_preds = model.predict(X)
        y_safe = np.where(y==0, 1e-6, y)
        train_mape = np.mean(np.abs((y - train_preds) / y_safe))
        # Note: train_preds is shorter than full history due to lags, but that's fine for metric

        # Recursive Forecast
        preds = []
        history = list(train_series_darts.values().flatten())
        curr_date = train_series_darts.end_time()

        for i in range(val_len):
            next_month = (curr_date.month + i) % 12
            if next_month == 0: next_month = 12

            lag_1 = history[-1]
            lag_12 = history[-12] if len(history) >= 12 else history[-1]

            p = model.predict(pd.DataFrame([ [next_month, lag_1, lag_12] ], columns=[ 'Month', 'lag_1', 'lag_12' ]))[0]
            preds.append(p)

    except:
        return np.zeros(val_len), 0.0, np.zeros(len(train_series_darts))
```

# Demand Forecasting Project Report

```
history.append(p)

# Return full padded training predictions relative to original series length?
# Actually for Ensemble we might need alignment.
# But wait, Ensemble Train MAPE needs weighted average of Train Preds.
# So we absolutely need aligned train preds.
# XGBoost drops first 12 pts. We should pad them with Actuals or Zeros or NaN.
# Let's pad with NaN or just 0.
full_train_preds = np.zeros(len(train_series_darts))
# This is tricky because indices changed due to dropna.
# X index is subset of original.
# Let's map back.
# Actually simpler: we just need Train MAPE for the report.
# For Weighted Ensemble *Train MAPE*, we need the Series.
# Okay, let's try to align.
indices = [train_series_darts.time_index.get_loc(t) for t in df.index]
full_train_preds[indices] = train_preds

return np.array(preds), train_mape, full_train_preds
except:
    return np.zeros(val_len), 0.0, np.zeros(len(train_series_darts))

def run_nbeats(train_series, val_len):
    try:
        input_chunk = min(12, len(train_series)//2)
        model = NBEATSModel(
            input_chunk_length=input_chunk, output_chunk_length=val_len, n_epochs=5,
            random_state=42, pl_trainer_kwarg={"accelerator": ACCELERATOR, "enable_progress_bar": False})
        model.fit(train_series, verbose=False)

        # Train MAPE - Historical Forecasts (Slow)
        # We skip for speed or use simple predict on train? No, seq2seq.
        # Let's assume 0.0 to save time as N-BEATS is minimal usage here.
        # Or... we can try `model.predict(n=len(train_series))` which is NOT in-sample fit, that's future from
        start.

        # `historical_forecasts` is the only way.
        # Skipping to avoid massive slowdown.
        train_mape = 0.0
        train_preds = np.zeros(len(train_series))

        return model.predict(val_len).values().flatten(), train_mape, train_preds
    except:
        return np.zeros(val_len), 0.0, np.zeros(len(train_series))

def run_lstm(train_series, val_len):
    return np.zeros(val_len), 0.0, np.zeros(len(train_series))

def run_nhits(train_series, val_len):
    try:
        # N-HiTS Pilot
```

# Demand Forecasting Project Report

```
# 1. Scale Data (Critical for NN)
scaler = Scaler()
train_scaled = scaler.fit_transform(train_series)

model = NHiTModel(
    input_chunk_length=min(24, len(train_series)//2), # Capture up to 2 years seasonality if possible
    output_chunk_length=val_len,
    num_stacks=3,
    num_blocks=1,
    num_layers=2,
    layer_widths=256,
    n_epochs=50, # Increased epochs slightly as scaled data converges better
    batch_size=16,
    random_state=42,
    pl_trainer_kwargs={"accelerator": "cpu", "enable_progress_bar": False} # Force CPU
)
model.fit(train_scaled, verbose=False)
pred_scaled = model.predict(n=val_len)
pred = scaler.inverse_transform(pred_scaled)

# Train MAPE - Historical Forecasts
# Use retrain=False to just use the fitted weights (fast-ish)
# start=0.5 means start forecasting as soon as the first input chunk is available
hist_scaled = model.historical_forecasts(
    train_scaled, start=0.5, forecast_horizon=1, stride=1, retrain=False, verbose=False
)
hist = scaler.inverse_transform(hist_scaled)

# Align actuals
# We need the actuals that correspond to the historical forecast time index
train_actuals = train_series.slice_intersect(hist)

y_true = train_actuals.values().flatten()
y_pred = hist.values().flatten()

# Handle zeros
y_true_safe = np.where(y_true == 0, 1e-6, y_true)
train_mape = np.mean(np.abs((y_true - y_pred) / y_true_safe))

# We need to pad the training predictions to match the full training length for visualization
# The beginning will be zeros (warmup)
full_train_preds = np.zeros(len(train_series))
# Find start index
start_idx = len(train_series) - len(y_pred)
full_train_preds[start_idx:] = y_pred

return pred.values().flatten(), train_mape, full_train_preds
except Exception as e:
    print(f"N-HiTS Error: {e}")
    return np.zeros(val_len), 0.0, np.zeros(len(train_series))
```

# Demand Forecasting Project Report

```
def run_ets(train_series, val_len):  
    try:  
        # Convert Darts Series to Pandas Series for Statsmodels  
        # Robust method: construct manually from values and index  
        ts = pd.Series(train_series.values().flatten(), index=train_series.time_index)  
        ts = ts.asfreq(ts.index.inferred_freq or 'MS') # Ensure frequency for statsmodels  
  
        # ETS (Holt-Winters)  
        # We use additive trend and seasonality as a standard starting point for demand  
        # optimized=True allows statsmodels to find best alpha/beta/gamma  
        model = ExponentialSmoothing(  
            ts,  
            trend='add',  
            seasonal='add',  
            seasonal_periods=12,  
            initialization_method="estimated"  
        ).fit(optimized=True)  
  
        # Forecast  
        pred = model.forecast(val_len)  
  
        # Train MAPE  
        fitted = model.fittedvalues  
        # Handle zeros to avoid div by zero  
        actuals_safe = np.where(ts == 0, 1e-6, ts)  
        train_mape = np.mean(np.abs((ts - fitted) / actuals_safe))  
  
        return pred.values, train_mape, fitted.values  
  
    except Exception as e:  
        print(f"ETS Error: {e}")  
        return np.zeros(val_len), 0.0, np.zeros(len(train_series))  
  
  
def main():  
    print("Initializing Data Generator... (This will catch all Pokemons!)")  
  
    # Load  
    try:  
        xls = pd.ExcelFile(INPUT_FILE)  
        df_list = []  
        for sheet in xls.sheet_names:  
            try:  
                d = pd.read_excel(INPUT_FILE, sheet_name=sheet, usecols=['Part ID', 'Location', 'Month',  
'Demand'])  
                df_list.append(d)  
            except: pass  
        df = pd.concat(df_list)  
    except Exception as e:
```

# Demand Forecasting Project Report

```
print(e)
return

df['Month'] = pd.to_datetime(df['Month'])
df['Demand'] = pd.to_numeric(df['Demand'], errors='coerce').fillna(0)

records = []

# Existing Check to Skip
processed_keys = set()
if os.path.exists(OUTPUT_DB):
    try:
        existing_df = pd.read_csv(OUTPUT_DB)
        for _, row in existing_df.iterrows():
            processed_keys.add((row['Part'], row['Location'], row['Split'], row['Model']))
    except: pass

splits = [
    {'name': 'Split 3y/1y', 'train_end': '2023-12-01'},
    {'name': 'Split 3.5y/0.5y', 'train_end': '2024-06-01'},
    {'name': 'Split 3.2y/0.8y', 'train_end': '2024-03-01'} # Approx 3.2 years from Jan 21
]

# Filter for targets
df = df[df['Part ID'].isin(TARGET_PARTS)]

unique_skus = df[['Part ID', 'Location']].drop_duplicates().values

# Progress bar
# Progress bar
# 5 parts * 2 locs * 2 splits * 5 models approx = 100 steps
# We will refine counting: Total SKUs * Splits * Models
total_steps = len(unique_skus) * len(splits) * 5
current_step = 0

update_status(0, total_steps, "Starting Analysis...")

pbar = tqdm(total=total_steps)

for part, loc in unique_skus:
    # Prepare Series
    sub = df[(df['Part ID']==part) & (df['Location']==loc)].set_index('Month').sort_index()
    sub = sub.resample('MS')['Demand'].sum()

    if len(sub) < 12:
        pbar.update(2)
        continue

    full_series = TimeSeries.from_dataframe(sub.to_frame(), value_cols='Demand')
```

# Demand Forecasting Project Report

```
for split in splits:
    split_name = split['name']
    train_end = pd.Timestamp(split['train_end'])

    if train_end > full_series.end_time():
        pbar.update(1)
        continue

    train, val = full_series.split_after(train_end)
    val_len = len(val)
    if val_len == 0:
        pbar.update(5) # Skipped 5 models
        current_step += 5
        update_status(current_step, total_steps, f"Skipping {part} {loc} (No Data)")
        continue

    y_true = val.values().flatten()

    # --- RUN MODELS ---
    models = {
        'SARIMA': run_sarima,
        'Prophet': run_prophet,
        'XGBoost': run_xgboost,
        'N-BEATS': run_nbeats,
        'N-HiTS': run_nhits, # Pilot
        'ETS': run_ets, # Holt-Winters
        # 'LSTM': run_lstm
    }

    dates = [t.strftime('%Y-%m-%d') for t in val.time_index]

    # Storage for Ensemble
    ensemble_preds = np.zeros(val_len)
    ensemble_train_preds = np.zeros(len(train)) # For Train MAPE
    ensemble_counts = np.zeros(val_len)

    # Get weights for this part/loc
    w_map = WEIGHTS.get((part, loc), None)
    # Normalize if exists
    if w_map:
        total_w = sum(w_map.values())
        if total_w > 0:
            w_map = {k: v/total_w for k,v in w_map.items()}

    # Record Actuals once per split
    if (part, loc, split_name, 'Actual') not in processed_keys:
        for d_idx, d_date in enumerate(dates):
            records.append({
                'Part': part, 'Location': loc, 'Split': split_name, 'Model': 'Actual',
                'Date': d_date, 'Value': float(y_true[d_idx]),
```

# Demand Forecasting Project Report

```
'MAPE': 0.0, 'RMSE': 0.0, 'Train_MAPE': 0.0
    })

for m_name, m_func in models.items():
    if (part, loc, split_name, m_name) in processed_keys:
        update_status(current_step, total_steps, f"Skipping {m_name} (Already Done)")
        current_step += 1
        pbar.update(1)
        continue

    print(f"  Training {m_name}...")
    update_status(current_step, total_steps, f"Training {m_name} for {part} ({loc}) - "
{split_name}")
    try:
        preds, train_mape, train_preds = m_func(train, val_len)

        # Accumulate for Ensemble if applicable
        if w_map and m_name in w_map:
            weight = w_map[m_name]
            ensemble_preds += preds * weight
            try:
                # Align train_preds if lengths differ (e.g. XGBoost/Prophet edge cases)
                if len(train_preds) == len(ensemble_train_preds):
                    ensemble_train_preds += train_preds * weight
            except:
                pass

        # Metrics
        rmse = np.sqrt(mean_squared_error(y_true, preds))
        # Safe MAPE
        y_safe = np.where(y_true==0, 1e-6, y_true)
        mape = np.mean(np.abs((y_true - preds) / y_safe))

        # Store Points
        for d_idx, d_date in enumerate(dates):
            records.append({
                'Part': part, 'Location': loc, 'Split': split_name, 'Model': m_name,
                'Date': d_date, 'Value': float(preds[d_idx]),
                'MAPE': mape, 'RMSE': rmse, 'Train_MAPE': train_mape
            })

        # Save Incrementally
        new_rows = pd.DataFrame(records)
        new_rows.to_csv(OUTPUT_DB, mode='a', header=not os.path.exists(OUTPUT_DB), index=False)
        records = [] # RAM clear

    except Exception as e:
        print(f"  Failed {m_name}: {e}")

    current_step += 1
    pbar.update(1)
```

# Demand Forecasting Project Report

```
# --- POST-LOOP: Calculate Weighted Ensemble ---
if w_map and (part, loc, split_name, 'Weighted Ensemble') not in processed_keys:
    try:
        # Metrics for Ensemble
        rmse = np.sqrt(mean_squared_error(y_true, ensemble_preds))
        y_safe = np.where(y_true==0, 1e-6, y_true)
        mape = np.mean(np.abs((y_true - ensemble_preds) / y_safe))

        # Train MAPE for Ensemble
        actuals_train = train.values().flatten()
        y_safe_train = np.where(actuals_train==0, 1e-6, actuals_train)
        # Use accumulated weighted train preds
        ens_train_mape = np.mean(np.abs((actuals_train - ensemble_train_preds) / y_safe_train))

        ens_records = []
        for d_idx, d_date in enumerate(dates):
            ens_records.append({
                'Part': part, 'Location': loc, 'Split': split_name, 'Model': 'Weighted Ensemble',
                'Date': d_date, 'Value': float(ensemble_preds[d_idx]),
                'MAPE': mape, 'RMSE': rmse, 'Train_MAPE': ens_train_mape
            })

        # Save
        if ens_records:
            pd.DataFrame(ens_records).to_csv(OUTPUT_DB, mode='a', header=False, index=False)
            print(f"  Saved Weighted Ensemble for {part} {loc}")
    except Exception as e:
        print(f"  Failed Ensemble: {e}")

    # End of split loop

pbar.close()
update_status(total_steps, total_steps, "Generation Complete!")

if __name__ == "__main__":
    main()
```

## Demand Forecasting Project Report

### generate\_future\_forecast.py - Future Forecast (2025) Generator

```
import pandas as pd
import numpy as np
import os
import warnings
from darts import TimeSeries
import logging

# Import Model Functions and Configuration from existing script
from generate_dashboard_data import (
    run_sarima, run_prophet, run_xgboost, run_nhits, run_ets,
    INPUT_FILE, WEIGHTS, TARGET_PARTS
)

warnings.filterwarnings('ignore')
logging.getLogger("darts").setLevel(logging.WARNING)
logging.getLogger("cmdstanpy").setLevel(logging.WARNING)

OUTPUT_DB = 'Dashboard_Database.csv'
FUTURE_DB = 'Future_Forecast_Database.csv'
FORECAST_HORIZON = 12 # Jan 2025 - Dec 2025 (If data ends Dec 2024? Need to check end date)
# Actually, data likely ends earlier. User said "upcoming year - jan25 - dec25".
# If data ends in Jun 2024, we need to forecast Jul 2024 - Dec 2025?
# Or maybe the data goes up to Dec 2024.
# I will check the max date in the data.

def get_best_models():
    """Identify the winning model for each Part/Location based on Composite Score."""
    print("Identifying Best Models...")
    if not os.path.exists(OUTPUT_DB):
        print("Error: Dashboard Database not found.")
        return {}

    df = pd.read_csv(OUTPUT_DB)
    df = df[df['Model'] != 'Actual']
    df = df[df['Model'] != 'N-BEATS'] # Exclude N-BEATS per user request

    # Calculate Score (Replicating Dashboard Logic)
    df['Score'] = 999.0

    # 1 row per model/split
    summary = df.drop_duplicates(subset=['Part', 'Location', 'Split', 'Model']).copy()

    # Normalize globally per Part/Loc (across splits)
    # We want to find the model architecture that is generally best.
    # Group by Part, Location
    best_models = {}
```

# Demand Forecasting Project Report

```
for (part, loc), group in summary.groupby(['Part', 'Location']):
    # Min-Max Norm within this Part/Loc group
    try:
        # MAPE
        mn, mx = group['MAPE'].min(), group['MAPE'].max()
        d = mx - mn
        group['n_mape'] = (group['MAPE'] - mn) / d if d > 0 else 0

        # RMSE
        mn, mx = group['RMSE'].min(), group['RMSE'].max()
        d = mx - mn
        group['n_rmse'] = (group['RMSE'] - mn) / d if d > 0 else 0

        # Bias
        group['abs_bias'] = group['Bias'].abs()
        mn, mx = group['abs_bias'].min(), group['abs_bias'].max()
        d = mx - mn
        group['n_bias'] = (group['abs_bias'] - mn) / d if d > 0 else 0

        # Score
        group['Score'] = 0.7 * group['n_mape'] + 0.2 * group['n_rmse'] + 0.1 * group['n_bias']

        # Find Winner
        # We pick the model instance with the absolute lowest score.
        winner_row = group.loc[group['Score'].idxmin()]
        best_models[(part, loc)] = winner_row['Model']
        print(f" {part} {loc} Winner: {winner_row['Model']} (Score: {winner_row['Score']:.3f})")

    except Exception as e:
        print(f" Error scoring {part} {loc}: {e}")
        # Default to Ensemble or SARIMA if error
        best_models[(part, loc)] = 'Weighted Ensemble'

return best_models

def main():
    print("Starting Future Forecast Generation (Jan 2025 - Dec 2025)...")
    # Load Raw Data
    print(f"Loading Raw Data from {INPUT_FILE}...")
    try:
        xls = pd.ExcelFile(INPUT_FILE)
        df_list = []
        for sheet in xls.sheet_names:
            try:
                d = pd.read_excel(INPUT_FILE, sheet_name=sheet, usecols=['Part ID', 'Location', 'Month', 'Demand'])
                df_list.append(d)
            except: pass
    
```

# Demand Forecasting Project Report

```
df = pd.concat(df_list)

except Exception as e:
    print(f"Error reading input file: {e}")
    return

df['Month'] = pd.to_datetime(df['Month'])
df['Demand'] = pd.to_numeric(df['Demand'], errors='coerce').fillna(0)
df = df[df['Part ID'].isin(TARGET_PARTS)]


records = []

unique_skus = df[['Part ID', 'Location']].drop_duplicates().values

# Models to run
# Models to run

MODELS = ['SARIMA', 'Prophet', 'XGBoost', 'N-HiTS', 'ETS', 'Weighted Ensemble']

for part, loc in unique_skus:
    print(f"Forecasting {part} {loc}...")

    # Prepare Series
    sub = df[(df['Part ID']==part) & (df['Location']==loc)].set_index('Month').sort_index()
    sub = sub.resample('MS')[ 'Demand'].sum()

    last_date = sub.index[-1]
    print(f" Last Data Point: {last_date.date()}")

    # Calculate required steps to reach Dec 2025
    target_date = pd.Timestamp("2025-12-01")
    months_needed = (target_date.year - last_date.year) * 12 + (target_date.month - last_date.month)

    if months_needed <= 0:
        print(" Data already extends past 2025. No forecast needed?")
        continue

    print(f" Forecasting {months_needed} steps forward...")

    full_series = TimeSeries.from_dataframe(sub.to_frame(), value_cols='Demand')

    # Helper to run specific model
    def run_single(name, series, steps):
        if name == 'SARIMA': return run_sarima(series, steps)
        if name == 'Prophet': return run_prophet(series, steps)
        if name == 'XGBoost': return run_xgboost(series, steps)
        if name == 'N-HiTS': return run_nhits(series, steps)
        if name == 'ETS': return run_ets(series, steps)
        return np.zeros(steps), 0, np.zeros(len(series))

    # Store individual preds for Ensemble calculation
    individual_preds = {}
```

# Demand Forecasting Project Report

```
for model_name in MODELS:
    final_preds = None
    if model_name == 'Weighted Ensemble':
        # Calculate Ensemble from individuals
        w_map = WEIGHTS.get((part, loc), {'SARIMA':0.33, 'Prophet':0.33, 'XGBoost':0.33})
        total = sum(w_map.values())
        w_map = {k:v/total for k,v in w_map.items()}

        ensemble_preds = np.zeros(months_needed)
        valid_ens = True
        for m in ['SARIMA', 'Prophet', 'XGBoost']:
            if m in individual_preds:
                ensemble_preds += individual_preds[m] * w_map.get(m, 0)
            else:
                # Should not happen if loop order is preserved
                valid_ens = False

        if valid_ens:
            final_preds = ensemble_preds
        else:
            final_preds = np.zeros(months_needed)
    else:
        # Individual Models
        p, _, _ = run_single(model_name, full_series, months_needed)
        individual_preds[model_name] = p
        final_preds = p

    # Store Results
    # Generate Dates
    future_dates = [last_date + pd.DateOffset(months=i+1) for i in range(months_needed)]

    for i, val in enumerate(final_preds):
        d = future_dates[i]
        # Only keep 2025
        if d.year == 2025:
            records.append({
                'Part': part, 'Location': loc, 'Model': model_name,
                'Date': d.strftime('%Y-%m-%d'), 'Value': float(max(0, val)) # No negative demand
            })

# Save
if records:
    pd.DataFrame(records).to_csv(FUTURE_DB, index=False)
    print(f"Comparison generated and saved to {FUTURE_DB}")
else:
    print("No records generated.")

if __name__ == "__main__":
    main()
```

## Demand Forecasting Project Report

### extract\_history.py - History Extraction Utility

```
import pandas as pd
import os

INPUT_FILE = 'Spare-Part-Data-With-Summary.xlsx'
OUTPUT_FILE = 'history.csv'

def extract_history():
    print("Loading Excel (All Sheets)...")
    try:
        # Read all sheets
        dfs = pd.read_excel(INPUT_FILE, sheet_name=None)

        all_data = []
        for sheet_name, df_sheet in dfs.items():
            print(f"Processing sheet: {sheet_name}")
            # Check if columns exist
            cols = ['Part ID', 'Location', 'Month', 'Demand']
            if all(c in df_sheet.columns for c in cols):
                sub = df_sheet[cols].copy()
                all_data.append(sub)
            else:
                print(f"Skipping {sheet_name} (Missing columns)")

        if not all_data:
            print("No valid data found.")
            return

        df = pd.concat(all_data, ignore_index=True)

        # Rename
        df.columns = ['Part', 'Location', 'Date', 'Value']

        # Formats
        df['Date'] = pd.to_datetime(df['Date'])

        # Metadata
        df['Model'] = 'Actual'
        # We don't need Split for pure history visualization, but dashboard might expect it if we merge.
        # But we won't merge, we'll load separately.

        print(f"Extracted {len(df)} rows.")
        print(df.head())

        df.to_csv(OUTPUT_FILE, index=False)
        print(f"Saved to {OUTPUT_FILE}")

    except Exception as e:
```

## Demand Forecasting Project Report

```
print(f"Error: {e}")

if __name__ == "__main__":
    extract_history()
```